



**SIGGRAPH**2015  
Xroads of Discovery

Real-time  
Managers  
and Tools



**SIGGRAPH 2015**  
Xroads of Discovery

The 42nd International Conference and Exhibition  
on Computer Graphics and Interactive Techniques



**CHALMERS**  
UNIVERSITY OF TECHNOLOGY

## Introduction to Real-Time Shading with Many Lights

Ola Olsson  
Chalmers University of Technology

- The first part then is going to be presented by me.
- In this part we will provide an overview of techniques that have been, and still are, used in real-time shading.
- The main focus will be on clustered shading, as this is the most advanced and efficient technique today...
- ...and is what the remaining talks in this course build on.

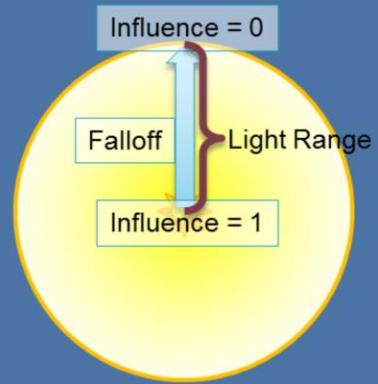
# Talk Outline

- Problem Definition
  - What, exactly, are you talking about!?!
- Part 1: 'Ancient' Shading Techniques
  - Forward Shading
  - Deferred Shading
- Part 2: Modern Shading Techniques
  - Tiled Deferred/Forward Shading
  - Clustered Deferred/Forward Shading

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

# Problem Definition

- Real-time shading algorithms
  - Thousands of lights
  - Limited range light
  - Fully dynamic
    - Lights and Geometry



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- The algorithms we are going to talk about are relevant when we target hundreds or thousands of lights in real time,
- And in a way you could say that the end goal is to make the number a non-issue. <click>
- The lights we consider have a limited range, with some falloff which goes to 0 at the boundary.
- This means that lights are not physical, but this is (still) the normal procedure in games. <click>
- We also do not consider shadows, which is covered later in the course.
- There is no pre-computation so all geometry and lights are allowed to change freely from frame to frame.

# Light Assignment – Not Shading!

- Question:
  - *Which lights affect what shading point.*
    - 2M pixels × 1k lights = enough
- Goals:
  - Minimize number of lights / pixel
  - High performance
- Solution:
  - *Light Assignment*
  - *(Light Culling)*

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Despite the name of all the techniques in this course, it is not actually about shading at all!
- This is a name we have inherited from forward shading / deferred shading etc.
- In reality the course is about working out which lights are relevant for shading each point of interest.
- With millions of sample points (pixels) and thousands of lights
- Doing this with high performance becomes sufficiently complex to be an interesting problem.
- We call this process ‘light assignment’, but is also known as ‘light culling’.

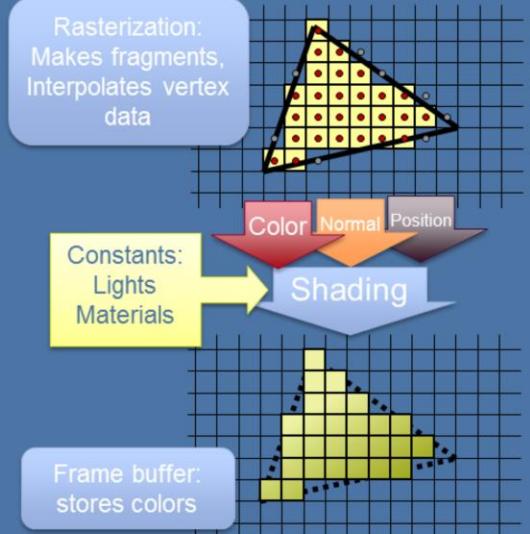
Part 1

# TRADITIONAL FORWARD SHADING

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

# (Traditional) Forward Shading

- The traditional approach
  - Single pass
- Rasterization generates fragments
  - Geometry samples
  - Interpolated attributes
- Each fragment is shaded
  - Needs set of lights.
  - Produces a color



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- The traditional method for real-time shading is called forward shading and used to be the only method during the first decade, or so, of consumer GPUs.
- It is still dominant on mobile hardware.
- In this technique, there is only a single pass over the geometry drawing into a frame buffer accumulating the final image.
- Geometry is rasterized, shading is performed and the frame buffer is updated.
- Shading is performed in the fragment shaders. (Or indeed, in vertex shaders...)

# Forward Shading

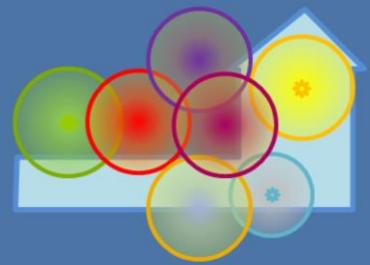
- Which lights to use?
  - Gather before draw call.
- Want minimal set of lights.
  - Small batches.
- Want quick drawing
  - Few batches!
    - Large batches.
  - Lower GPU/API overhead.
  - Fewer material changes.

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Thus, we need to round up the set of relevant lights before each draw call.
- and assign lights per chunk/batch of primitives.
- To minimize number of lights, we want to make sure a batch is small, geometrically.
- But to draw fast, keep the GPU busy and avoid API overhear, we want large batches.
- And don't care so much about geometric shape.
- This is a fundamental conflict, where the best we can manage is a compromise. But it is a difficult one to strike.

# Forward Shading Problems

- Few large objects
- Many small lights

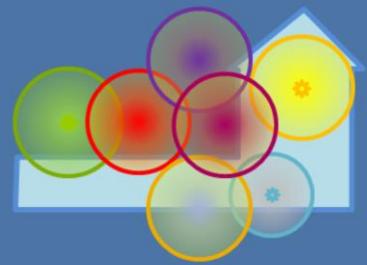


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Now, this conflict runs deeper than simply batch size.
- The basic problem is that we have to assign lights based on the size of geometry chunks.
- Therefore, on the one hand, we might have a situation with a few large objects, and many small lights

# Forward Shading - Problems

- Few large objects
- Many small lights
- Redundant shading work

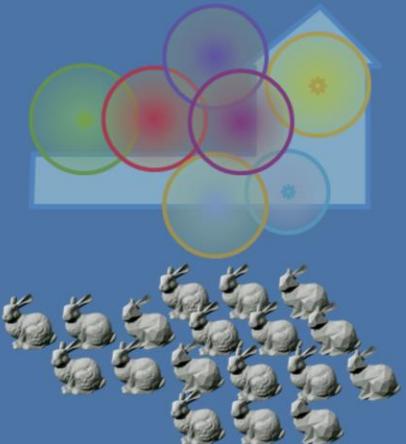


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- This leads to lots of wasted effort in shading lights that only affect a portion of the geometry.

# Forward Shading - Problems

- Few large objects
- Many small lights
- Redundant shading work
- Many small objects

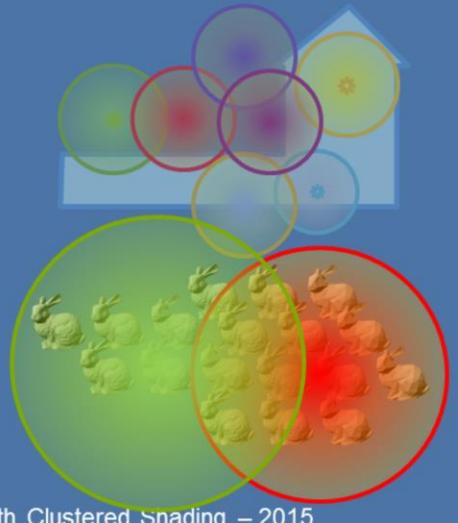


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- On the other hand, we may have many small objects.

# Forward Shading - Problem Examples

- Few large objects
- Many small lights
- Redundant shading work
- Many small objects
- Few large lights

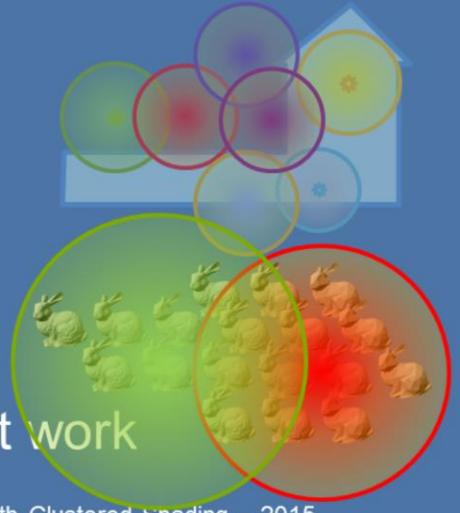


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- And a few large lights, <click>

# Forward Shading - Problem Examples

- Few large objects
- Many small lights
- Redundant shading work
- Many small objects
- Few large lights
- Redundant light assignment work

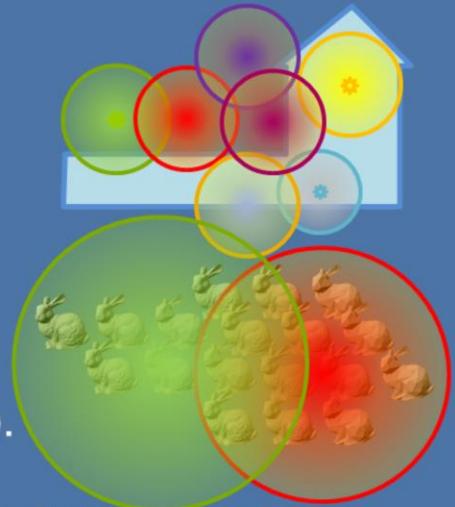


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- in which case we will spend a lot of effort to individually discover that they are affected by the same lights.

# Forward Shading - Problem Examples

- Few large objects
- Many small lights  
AND
- Many small objects
- Few large lights
  
- Large triangles.
- Varying light density (dynamic).



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- So the conflict runs deeper than just batch size...<click>
- And of course, in the same scene we might have both houses and rabbits, which means it is difficult or impossible to ensure a good performance balance.
- Note that the large object could be a single triangle, and in general triangle size and varying light density makes this a very difficult problem to solve well.

# Summary: Forward Shading

## The good

- Single Pass
- Low storage overhead
  - Single Frame Buffer
- Transparency works.
- MSAA works
- Simple if only few lights
  - E.g., the sun
- Varying shading models is easy.

## The Bad

- Overdraw
- Batching coupled with lighting
- Shader management
  - Permutations of #lights/type.

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- To summarize, the main good thing about traditional forward shading is that

# Summary: Forward Shading

## The good

- Single Pass
- Low storage overhead
  - Single Frame Buffer
- Transparency works.
- MSAA works
- Simple if only few lights
  - E.g., the sun
- Varying shading models is easy.

## The Bad

- Overdraw
- Batching coupled with lighting
- Shader management
  - Permutations of #lights/type.

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Everything can be done in a single pass with a single frame buffer, just storing the final color
- This enables supporting transparency and MSAA without much ado.

# Summary: Forward Shading

## The good

- Single Pass
- Low storage overhead
  - Single Frame Buffer
- Transparency works.
- MSAA works
- Simple if only few lights
  - E.g., the sun
- Varying shading models is easy.

## The Bad

- Overdraw
- Batching coupled with lighting
- Shader management
  - Permutations of #lights/type.

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Finally, we have that problem of assigning lights at the right granularity which makes an efficient implementation difficult.

# TRADITIONAL DEFERRED SHADING

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

# (Traditional) Deferred Shading

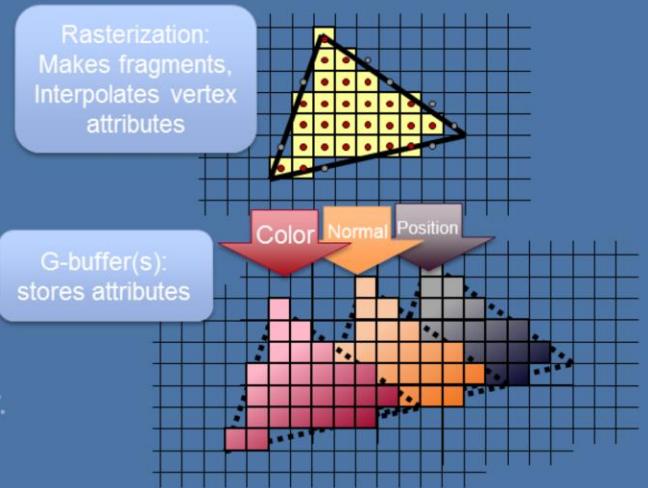
- Goal:
  - Decouple shading from geometric complexity
  - Solve overdraw (shade once/sample)
- Modern form introduced in 1990 [Saito90]
  - Term coined later by Tebbs et al [Tebbs92]
  - Consumer GPUs in 2004 [Hargreaves04]
  - Games around 2005 [Shishkovtsov05]

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- The basic goal is to solve both the overdraw problem and light assignment problem by decoupling geometry from shading.<click>
- Deferred shading as we know it today with G-Buffers was introduced in 1990.
- Although the term ‘deferred shading’ actually was coined later.
- Started appearing in games around 2005, with the introduction of MRT.

# Deferred Shading

- Key idea
  - Defer shading
- First do a geometry pass
  - Sample geometry
  - Interpolate attributes
  - But, no shading
  - Instead: G-buffers
- Then do a shading pass
  - Process lights independently.
  - Shade only visible samples.

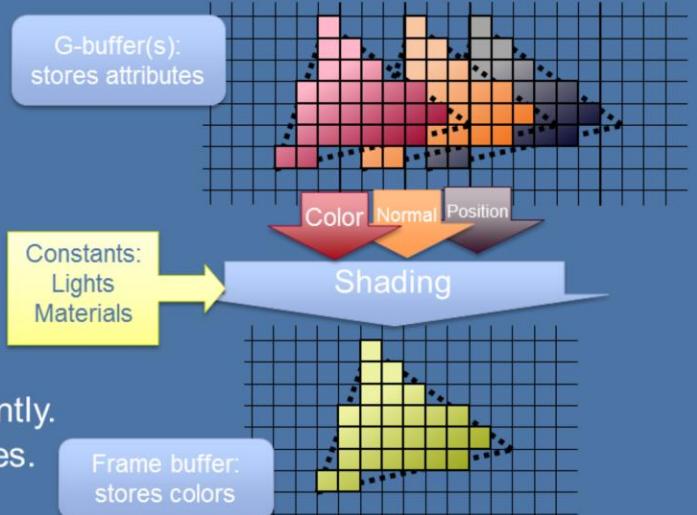


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- The key idea is to defer, that is wait with, the shading to its own pass.
- So in a first pass, triangle geometry is transformed and sampled into fragments.
- But instead of shading, the per-pixel geometry attributes are stored in what is called G-buffers.
- The important thing is that they store all the attributes that can change per pixel.<click>

# Deferred Shading

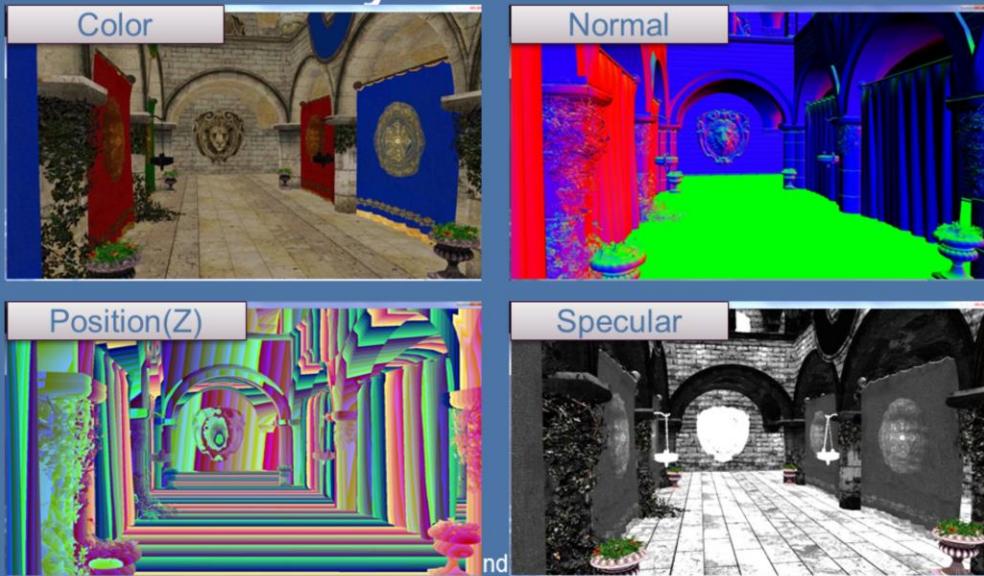
- Key idea
  - Defer shading
- First do a geometry pass
  - Sample geometry
  - Interpolate attributes
  - But, no shading
  - Instead: G-buffers
- Then do a shading pass
  - Process lights independently.
  - Shade only visible samples.



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- When all geometry has been processed a separate shading pass is performed.
- And this pass is then independent of the scene geometry representation, and only works on the actually visible samples.

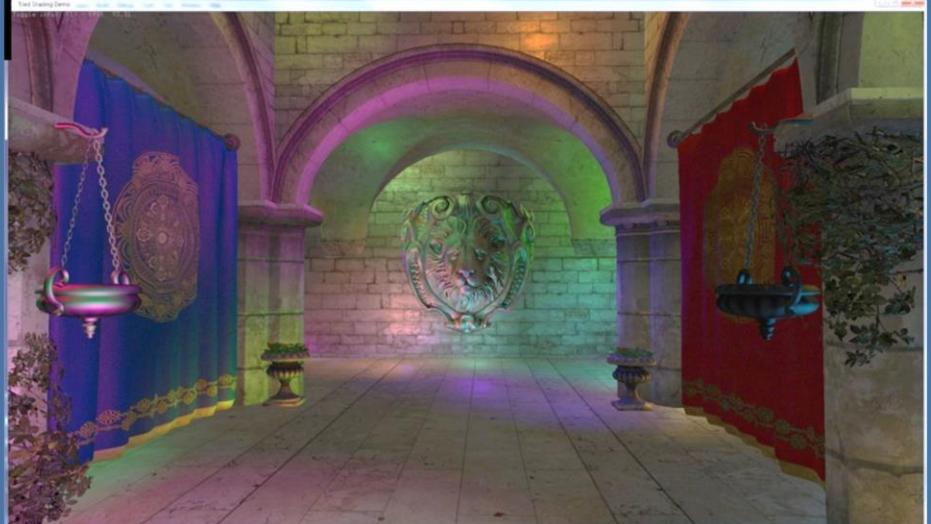
# Geometry Pass: G-Buffers



- So a set of G-Buffers might look like this.

# Shading Pass

- For each light
  - Draw Light Bounds
- For each fragment
  - Read G-Buffers
  - Compute Shading
  - Add to frame buffer



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- The way we solve the light assignment is to use the rasterizer to draw the bounding volumes of each light <click> as geometry to the screen<click>
- In the fragment shader of these we compute the shading.
- Which is then blended into the frame buffer.
- This is done independently for each light.

# Deferred Light Shader

```

uniform vec3 lightPosition;
uniform vec3 lightColor;
uniform float lightRange;

void main()
{
    vec3 color = texelFetch(colorTex, gl_FragCoord.xy);
    vec3 specular= texelFetch(specularTex, gl_FragCoord.xy);
    vec3 normal = texelFetch(normalTex, gl_FragCoord.xy);
    vec3 position = fetchPosition(gl_FragCoord.xy);

    vec3 shading = doLight(position, normal, color,
                           specular, lightPosition,
                           lightColor, lightRange);

    resultColor = vec4(shading, 1.0);
}
  
```



Real-Time Many-Light Management and Shadows with Clusters

15

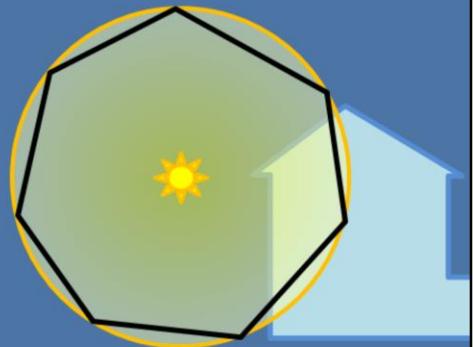
- Here is a somewhat simplified light shader.
- So this is a fragment shader that is used when drawing a *light*.
- It is thusly executed once for each pixel covered by the light bounding sphere.
- The shader has as uniform parameters the light attributes, position etc, these are the same for each fragment.
- First all attributes are fetched from the G-Buffers.
  - Note that the screen space coordinate of the fragment is in `gl_FragCoord.xy`
- Then shading is computed using these and the uniform attributes of the light.
- And output to the color of the pixel.

# Stencil Buffer Optimization

## [Arvo03]

- Analogous to stencil shadows.
- Draw light volume twice.
  - First back faces.
    - Set stencil mask on depth fail.
  - Then front faces.
    - Test stencil mask.
    - Compute shading.

View Direction

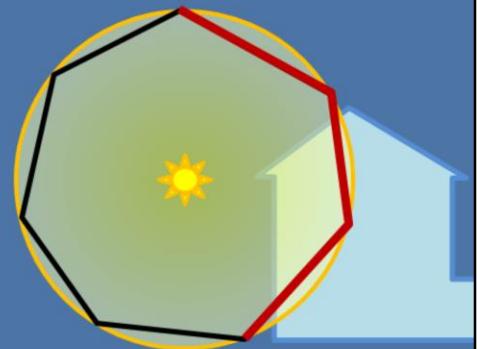


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

# Stencil Buffer Optimization

## [Arvo03]

- Analogous to stencil shadows.
- Draw light volume twice.
  - **First back faces.**
    - Set stencil mask on depth fail.
  - Then front faces.
    - Test stencil mask.
    - Compute shading.

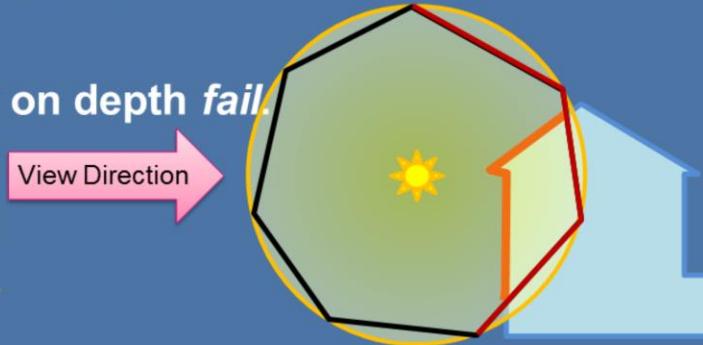


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

# Stencil Buffer Optimization

## [Arvo03]

- Analogous to stencil shadows.
- Draw light volume twice.
  - First back faces.
    - Set stencil mask on depth fail.
  - Then front faces.
    - Test stencil mask.
    - Compute shading.

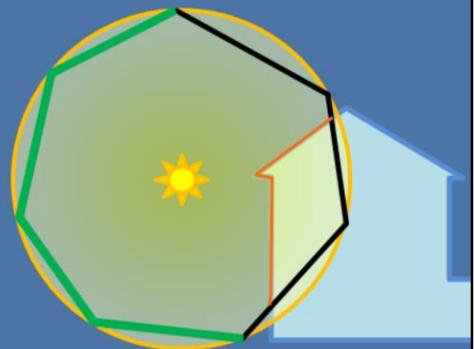


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

# Stencil Buffer Optimization

## [Arvo03]

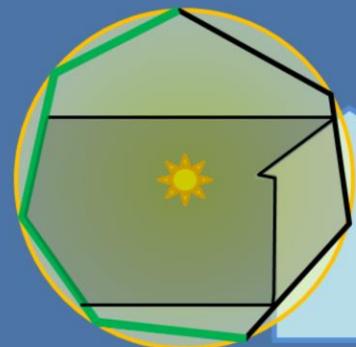
- Analogous to stencil shadows.
- Draw light volume twice.
  - First back faces.
    - Set stencil mask on depth fail.
  - **Then front faces.**
    - Test stencil mask.
    - Compute shading.



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

# Stencil Buffer Optimization [Arvo03]

- Analogous to stencil shadows.
- Draw light volume twice.
  - First back faces.
    - Set stencil mask on depth fail.
  - Then front faces.
    - **Test stencil mask.**
    - **Compute shading.**



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- A very important optimization is using the stencil buffer to exclude any pixels that are outside the light volume.
- This improves efficiency by only shading exactly those pixels that are needed.
- It is however unclear if it can be efficiently implemented on modern hardware for many lights.

# Other Variants

- Deferred Lighting
  - Factor out specular and diffuse color
  - G-Buffer only store normal and shininess
  - Output diffuse and specular shading
  - Second geometry pass which multiplies colors
- Light PrePass
  - Much like the above
  - But with monochromatic specular highlight
- Similar performance as deferred
  - Only improves constant factors.
- Limits shading model even further

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

# Summary: Trad. Deferred Shading

## The good

- Trivial light management
  - Enables many lights
- Simple (light) shader management
- High Shading efficiency
  - Only shade lit samples
- No overdraw
- Single geometry pass
- Shadow map reuse

## The Bad

- No transparency
- Large frame buffer
  - Especially with MSAA
  - Accumulates light in frame buffer
    - High precision needed
- Difficult to do multiple shading models
  - Custom shaders
- High memory bandwidth usage

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- The large advantages are, especially when the technique was introduced,
- Trivial light management, you just cull and draw the light bounds,
- And light shaders only need to support one light type.

# Summary: Trad. Deferred Shading

## The good

- Trivial light management
  - Enables many lights
- Simple (light) shader management
- High Shading efficiency
  - Only shade lit samples
- No overdraw
- Single geometry pass
- Shadow map reuse

## The Bad

- No transparency
- Large frame buffer
  - Especially with MSAA
  - Accumulates light in frame buffer
    - High precision needed
- Difficult to do multiple shading models
  - Custom shaders
- High memory bandwidth usage

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- The large advantages are
- Efficiency is very high, in that only samples that are actually affected by lights need to be shaded.

# Summary: Trad. Deferred Shading

## The good

- Trivial light management
  - Enables many lights
- Simple (light) shader management
- High Shading efficiency
  - Only shade lit samples
- No overdraw
- Single geometry pass
- Shadow map reuse

## The Bad

- No transparency
- Large frame buffer
  - Especially with MSAA
  - Accumulates light in frame buffer
    - High precision needed
- Difficult to do multiple shading models
  - Custom shaders
- High memory bandwidth usage

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- On the down side,
- Transparency is not readily supported

# Summary: Trad. Deferred Shading

## The good

- Trivial light management
  - Enables many lights
- Simple (light) shader management
- High Shading efficiency
  - Only shade lit samples
- No overdraw
- Single geometry pass
- Shadow map reuse

## The Bad

- No transparency
- Large frame buffer
  - Especially with MSAA
  - Accumulates light in frame buffer
    - High precision needed
- Difficult to do multiple shading models
  - Custom shaders
- High memory bandwidth usage

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- On the down side,
- We get very large frame buffers, especially with MSAA
- And because we need higher precision for light accumulation.

# Summary: Trad. Deferred Shading

## The good

- Trivial light management
  - Enables many lights
- Simple (light) shader management
- High Shading efficiency
  - Only shade lit samples
- No overdraw
- Single geometry pass
- Shadow map reuse

## The Bad

- No transparency
- Large frame buffer
  - Especially with MSAA
  - Accumulates light in frame buffer
    - High precision needed
- Difficult to do multiple shading models
  - Custom shaders
- High memory bandwidth usage

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- On the down side,
- It is more difficult to do custom shaders for certain geometry.

# Summary: Trad. Deferred Shading

## The good

- Trivial light management
  - Enables many lights
- Simple (light) shader management
- High Shading efficiency
  - Only shade lit samples
- No overdraw
- Single geometry pass
- Shadow map reuse

## The Bad

- No transparency
  - Large frame buffer
    - Especially with MSAA
    - Accumulates light in frame buffer
      - High precision needed
  - Difficult to do multiple shading models
- Custom shaders
- High memory bandwidth usage

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- There are many other pros and cons, but we'll focus on the most problematic:
  - High memory bandwidth requirements,

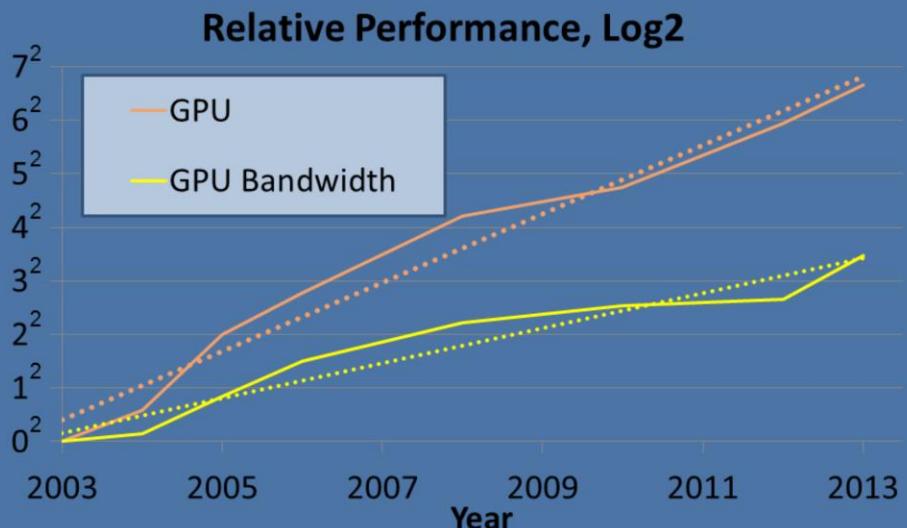
Part 2

# MODERN SHADING TECHNIQUES

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- This last problem brings us to the topic of modern shading techniques, so we'll look into this in some more detail.

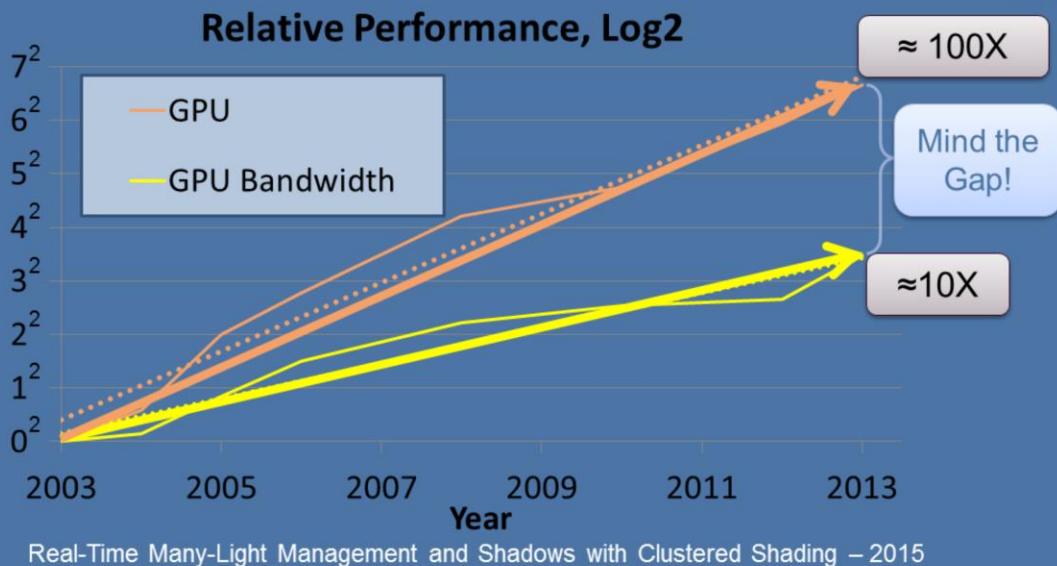
# GPU Performance Trends



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- The graph shows the relative performance trends of top NVIDIA GPUs over 10 years
- Showing clear exponential growth in both compute and bandwidth...

# GPU Performance Trends



- Note, however, the difference in exponent!<click>
- Over the 10 years plotted that amounts to a difference of a factor 10.<click>...
- ...for a compute versus memory bound algorithm.
- So, aiming for compute bound is key in future proofing your algorithms.

# Modern Shading Techniques



- Observations:
  1. GPU Compute >> Bandwidth
  2. Better GPU general purpose capability
- Conclusion:
  - Explore alternatives to rasterization!
    - When it makes sense...
- Broad impact
  - Many algorithms forced triangles!
    - Shadow volumes
    - Various splatting algorithms

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- With this, modern approaches to real-time shading with many lights all take their in the following 2 observations
  - First, as we just saw compute is king.
  - Second, the GPU general purpose programming models and power of these cores has improved tremendously over the last years.
- These observations have broad impact because not long ago triangle rasterization was the only game in town if you wanted real-time performance.
- This means that we are now sometimes stuck with a way of doing things that is no longer suitable.
- Thus I think the field is open to revisit many real-time algorithms with this frame of mind.

# The Bandwidth problem

- New type of overdraw
  - Light overdraw
- N lights cover a certain pixel ->
  - N reads from the **same** G-Buffer location

```

for each light
  for each covered pixel
    read G-Buffer
    compute shading
    read + write frame buffer
  
```

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Tiled and Clustered shading are examples of this, and it brings us back to the bandwidth problem of traditional deferred.
- As we are now drawing the lights, in the shading pass, we have overdraw when many lights overlap the same pixel.
- Schematically deferred shading looks like this, we iterate over each light,
- and then in parallel, by drawing the bounding volume, over all the fragments.

# The Bandwidth problem

- Repeated reads and writes

```
for each light
  for each covered pixel
    read G-Buffer
    compute shading
    read + write frame buffer
```

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- The problem is from the fact that the innermost loop is over the pixels,
- This which requires repeated reading, and writing of the G-Buffers and frame buffer.

# The Bandwidth problem

- Re-write loop!
- Hoist read/write outside.

```
for each light
  for each covered pixel
    read G-Buffer
    compute shading
    read + write frame buffer
```

Re-order loops  
Load/store -> Outer loop

```
for each pixel
  read G-Buffer
  for each affecting light
    compute shading
  write frame buffer
```

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- So we want to re-arrange this loop to make the innermost loop iterate over the lights instead. .<click>
- Then we get this nice compute oriented loop.
- With a single read and write which would solve the bandwidth problem.

# The Bandwidth problem

- Requires
  - sequential access to lights
- Global list?
  - Inefficient!
- List per pixel? [Trebilco07]
  - Slow construction  $\Leftrightarrow$  Shading
  - Lots of data

```

for each pixel
  read G-Buffer
  for each affecting light
    compute shading
  write frame buffer
  
```

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- However, we now need to access all the relevant lights for each pixel sequentially.
- Just using a global list of lights is of course terribly inefficient.
- At the other end of the spectrum, creating lists of lights for each pixel individually is both slow and requires lots of storage.

# The Modern shading solution

- Share between groups of *similar* pixels
  - Lots of coherency between samples
  - Coherent access
  - Little storage
  - Conservative lists

```

for each pixel
  read G-Buffer
  for each possibly affecting light
    if affecting
      compute shading
    write frame buffer
  
```

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- The modern solution is instead to build list for groups of similar samples.
- The motivation is of course that many pixels will use the same lights in a typical scenario
- Allowing us to share the list over many samples.
- Then list must be conservative, storing all lights that *may* affect any sample within the tile.
- So we trade some compute performance for bandwidth,
- which as we have seen is a pretty good gambit on modern GPUs.
- But at the same time significantly closer to the optimum compared to brute force.

# Modern Deferred Shading

- 1. Render Scene to G-Buffers
- 2. Light assignment
  - Build Light Acceleration Structure (Grid)
- 3. Full Screen Pass
  - Quad (or CUDA, or Compute Shaders, or SPUs)
  - For each pixel
    - Fetch G-Buffer Data
    - Look up light list in acceleration structure.
    - Loop over lights and accumulate shading
    - Write shading

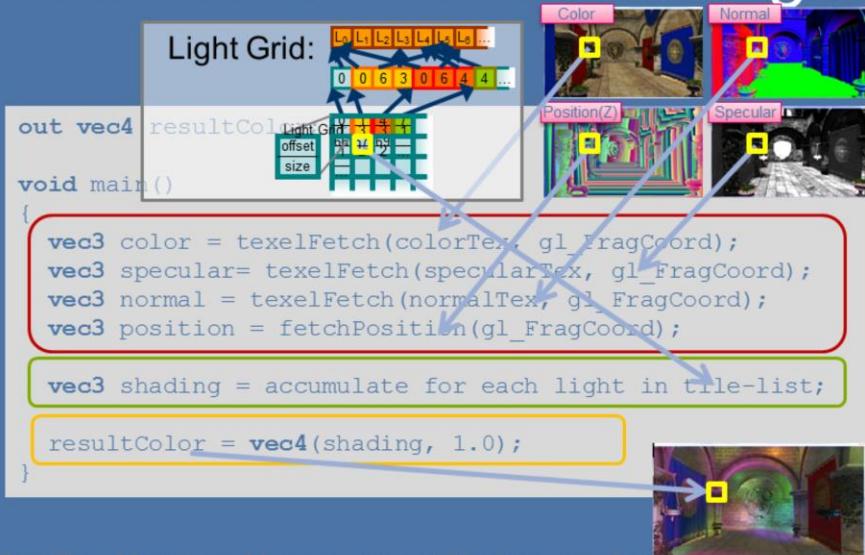
Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

# Modern Forward Shading

- (1. Optional: Pre-Z/Geometry Pass)
- 2. Light assignment
  - Build Light Acceleration Structure (Grid)
- 3. Geometry Pass
  - Just your normal shading pass.
  - For each fragment
    - Look up light list in acceleration structure.
    - Loop over lights and accumulate shading
    - Write shading to frame buffer

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

# Full-Screen Modern Shading Pass



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

# Summary: Modern Shading

## The good

- Low bandwidth
- Flexible
  - Forward or Deferred
- Transparency
- MSAA

## The Bad

- Complex light shader
  - If multiple light types
- No Shadow map reuse

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- So, we have pretty much solved the bandwidth problem which is great,
- It matters a lot less suddenly if the G-Buffers are fat as chips

# Summary: Modern Shading

## The good

- Low bandwidth
- Flexible
  - Forward or Deferred
- Transparency
- MSAA

## The Bad

- Complex light shader
  - If multiple light types
- No Shadow map reuse

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- The techniques are flexible, in that it is fairly trivial to support forward or deferred shading, or both.

# Summary: Modern Shading

## The good

- Low bandwidth
- Flexible
  - Forward or Deferred
- Transparency
- MSAA

## The Bad

- Complex light shader
  - If multiple light types
- No Shadow map reuse

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Using forward, in turn, makes it possible to support transparent and MSAA.

# Summary: Modern Shading

## The good

- Low bandwidth
- Flexible
  - Forward or Deferred
- Transparency
- MSAA

## The Bad

- Complex light shader
  - If multiple light types
- No Shadow map reuse

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- On the down side, the light shader now must support all light types, which makes it more complex.
- And since lights are not sequentially processed we cannot re-use shadow map storage space between lights.
  - More on this in my second talk!

Tiled Deferred, Tiled Forward (Forward+)

# TILED SHADING

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- This brings us to the first of the modern techniques that has been developed in this new, bandwidth constrained and compute oriented, landscape.
- Collectively called Tiled Shading, it covers both deferred and forward variants,
- The forward variant is also referred to as “Forward+” by AMD. I highly recommend using Tiled Forward Shading as this name really communicates something, for roughly the same reason that we don’t name our functions ‘foo’ and ‘bar’... right?

# Example Scene



- Crytek Sponza
  - Minus top
- 6 Lights
- A Tree

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- We will be using this example scene to illustrate how tiled and clustered shading works.
- The scene is the usual Crytek Sponza scene, but with the top three quarters lifted off, to let us get a better view.
- I've added six lights of different colours, represented as spheres.
- To make the view more interesting, I've also added a tree.

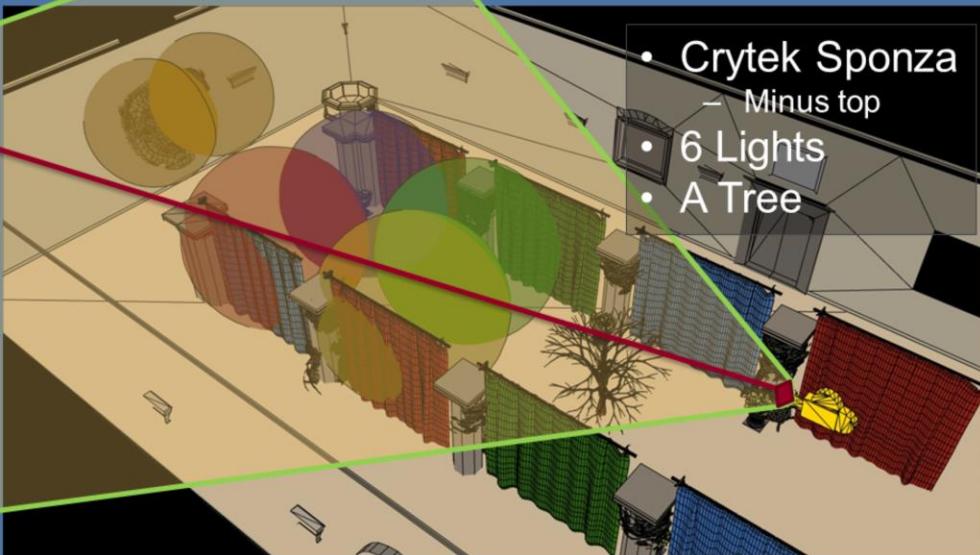
# Example Scene



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- The same view in outline mode, to make it more clear.

# Example Scene



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- This is the viewpoint that will be used to demonstrate the algorithms, the camera is looking through the tree towards the lion head on the wall.

# Example Scene

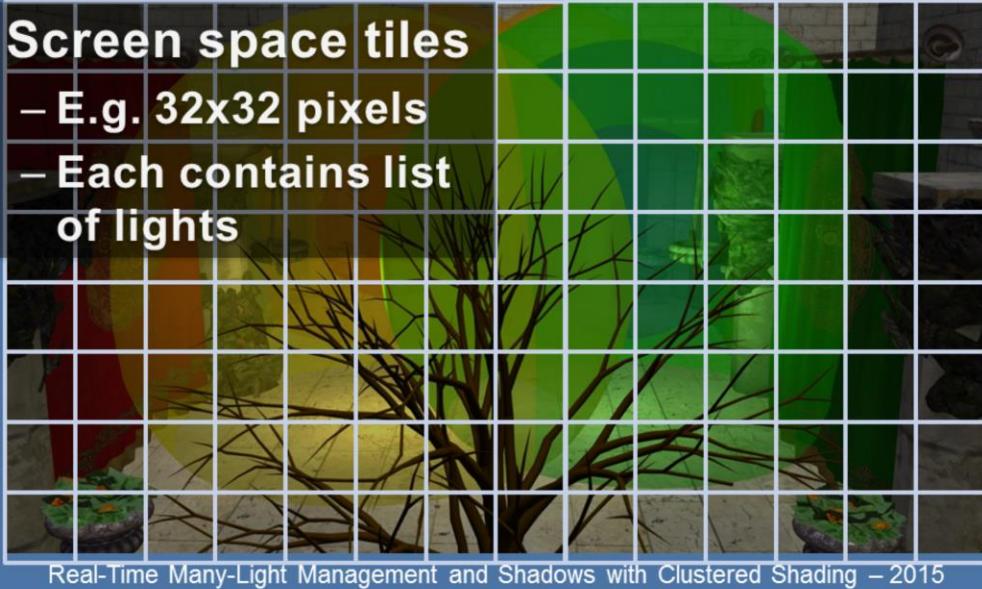


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Shifting to the point of view of the yellow camera.

# Tiled Shading Grid Construction

- **Screen space tiles**
  - E.g. 32x32 pixels
  - Each contains list of lights

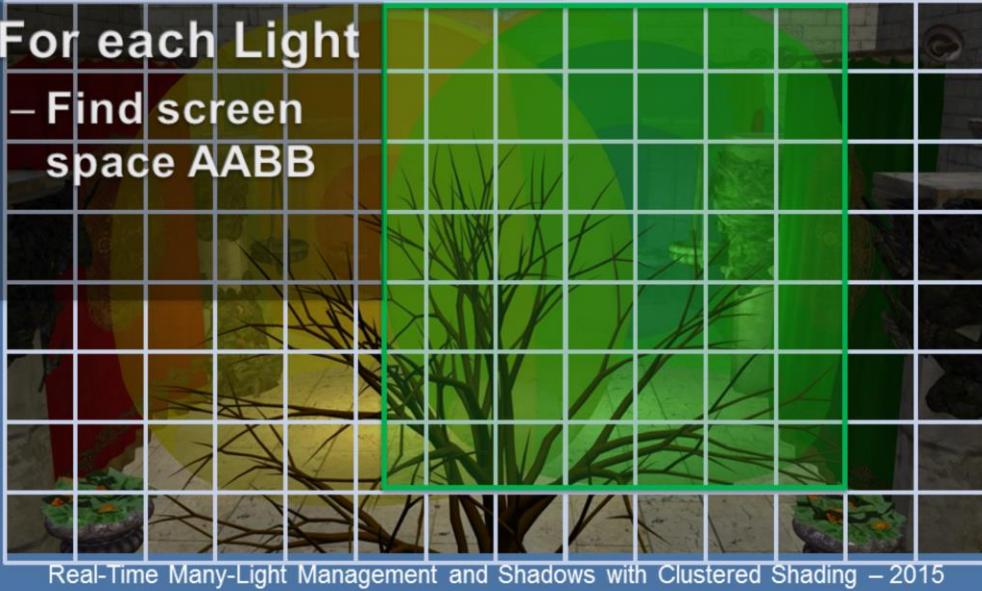


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- The screen is divided into tiles, each covering say 32 by 32 pixels.
- Each tile contains a single list of all the lights that might influence any of the pixels inside.
- Note that this list is shared between the pixels, so overhead for list maintenance and fetching is low.

# Tiled Shading Grid Construction

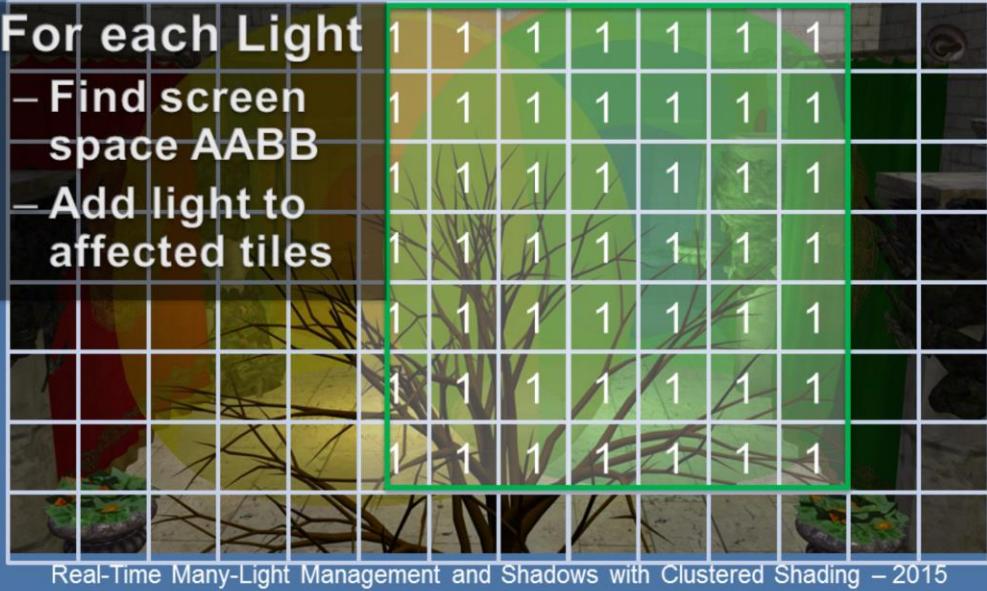
- For each Light
  - Find screen space AABB



- To construct the lists, we might do as follows.
- For each light, establish the screen space bounding box, illustrated for the green light.

# Tiled Shading Grid Construction

- **For each Light**
  - Find screen space AABB
  - Add light to affected tiles



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Then add the index of the light to all of the overlapped tiles.

# Tiled Shading Grid Construction



- Then repeat this process for all remaining lights.

# Tiled Shading Grid Construction

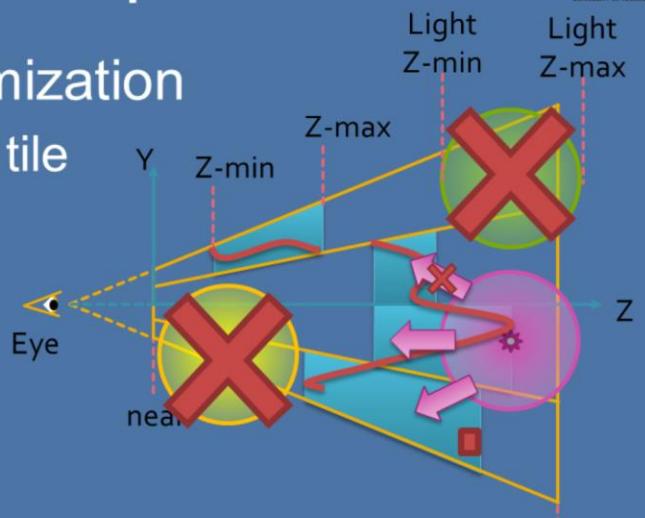


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- The illustration only shows the counts, so you need to imagine the lists being built as well.

# Depth Range Optimization

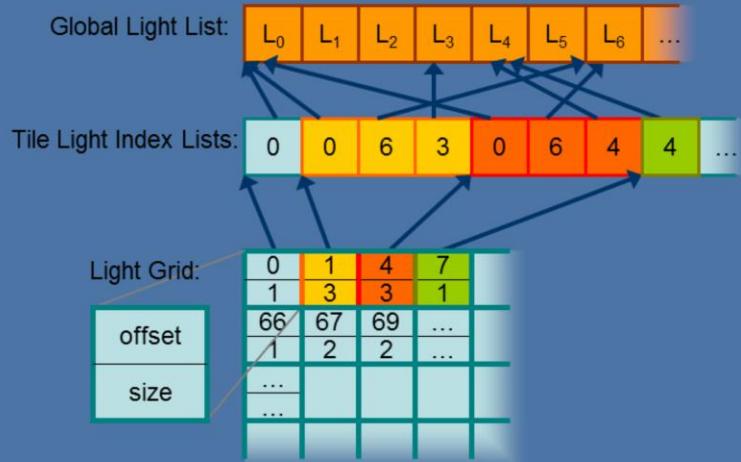
- Depth Range Optimization
  - Min/Max depth per tile
  - From Pre-z pass
- For each Light
  - Test depth bounds



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- In practice we'd also do a conservative per-tile min/max depth test, to cull away lights occupying empty space.
  - Just like the stencil optimization for traditional deferred.
- If we don't do this performance generally goes down the toilet...
- and this is what makes a pre-z pass necessary for tiled forward shading.

# Tiled Shading Light Grid



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- After we have processed all the lights, we end up with a 2D grid such as this
- This can then be used in the shading pass, as we saw before.

# Summary: Tiled Shading

## The good

- Low bandwidth
- Simple light assignment
- Trivial Light list lookup
- High performance
- Flexible
  - Forward or Deferred
- Transparency
- MSAA

## The Bad

- Complex light shader
- No Shadow map reuse
- View dependence
  - 2D light assignment
  - Depth discontinuities

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- To reiterate, we've solved the bandwidth problem

# Summary: Tiled Shading

## The good

- Low bandwidth
- Simple light assignment
- Trivial Light list lookup
- High performance
- Flexible
  - Forward or Deferred
- Transparency
- MSAA

## The Bad

- Complex light shader
- No Shadow map reuse
- View dependence
  - 2D light assignment
  - Depth discontinuities

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Light assignment is quite simple, just a 2D operation and lookup just depends on screen space location!

# Summary: Tiled Shading

## The good

- Low bandwidth
- Simple light assignment
- Trivial Light list lookup
- High performance
- Flexible
  - Forward or Deferred
- Transparency
- MSAA

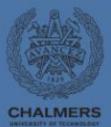
## The Bad

- Complex light shader
- No Shadow map reuse
- View dependence
  - 2D light assignment
  - Depth discontinuities

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

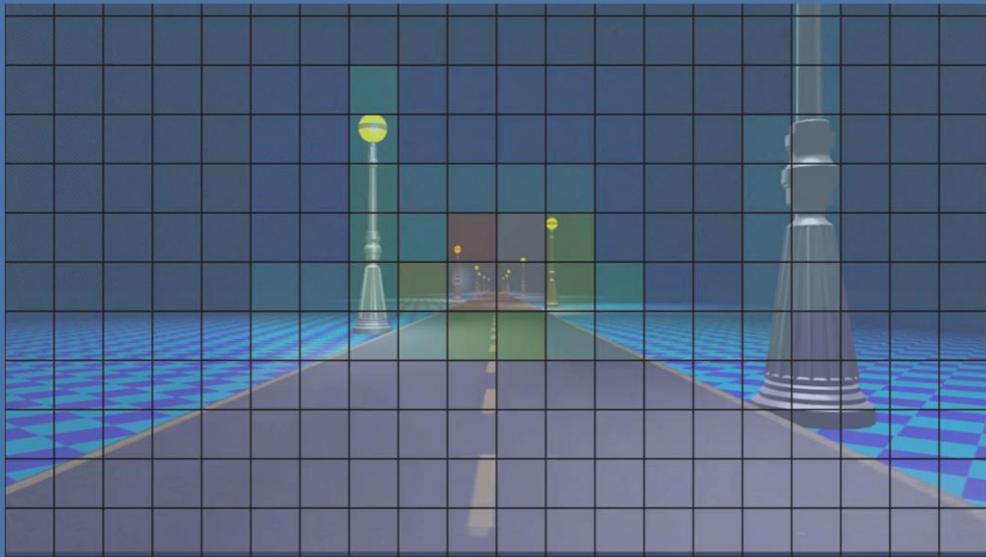
- However, performance is strongly view dependent.
- Which means run time performance is not readily predictable from scene parameters such as light distribution.
- We'll look into this problem in some detail next...

# Troublesome Tiles



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

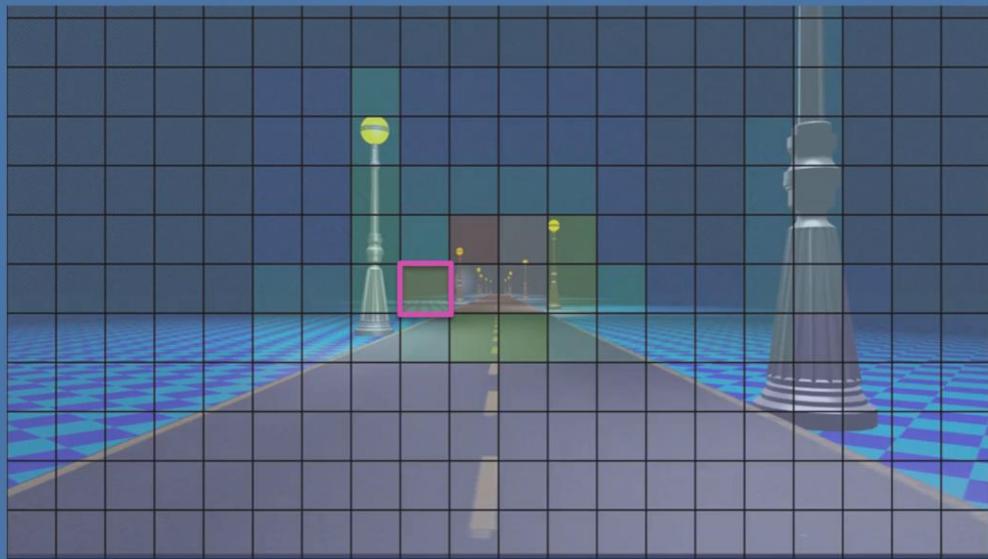
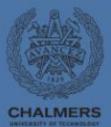
# Troublesome Tiles



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- So as we can see most of the trouble is around the light posts, where there are depth discontinuities.
- And roads aren't that uncommon in games I think.

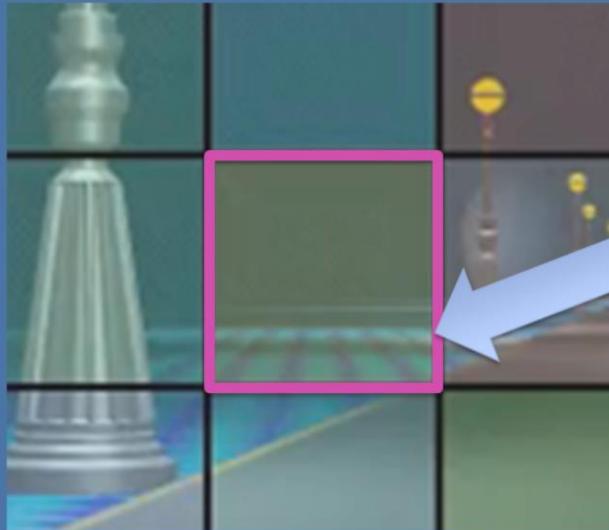
# Troublesome Tiles



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

\* But, as if that is not bad enough, lest focus on this tile.

# Troublesome Tiles



Flat Ground  
**No** discontinuities!

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Zoomed in we see there is only flat ground, and no discontinuities.
- So that should be a nice and good tile, right?

# Troublesome Tiles



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Well, no! It is pretty darn long.
- And I believe ground is something many games contain, so this might crop up here or there.

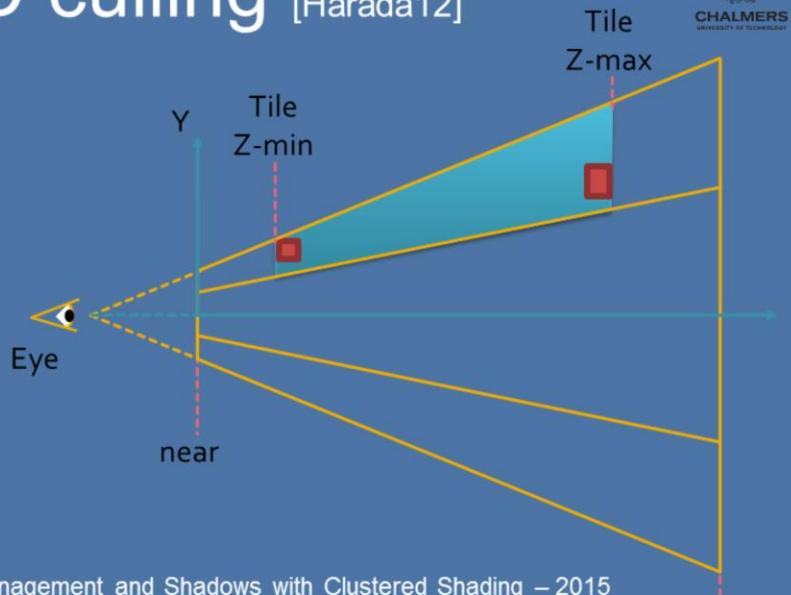
# Tiled extensions

- Attempts to address this problem
- 2.5D culling [Harada12]
- 'bimodal clusters' / 'HalfZ' [Lauritzen12][Thomas15]

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Some techniques that mitigate this problem exists
- I had to cull the slides describing how they work, but you can check out the GDC presentation by Gareth Thomas, where they are explained very well.

# 2.5D culling [Harada12]

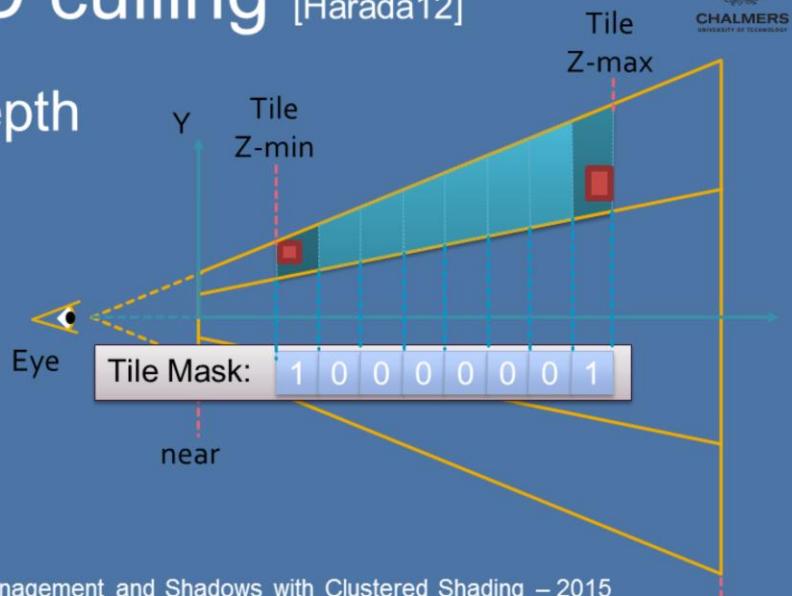


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Some techniques that attempts to address this exists
- One is 2.5D culling, which

# 2.5D culling [Harada12]

- Bit mask in depth direction

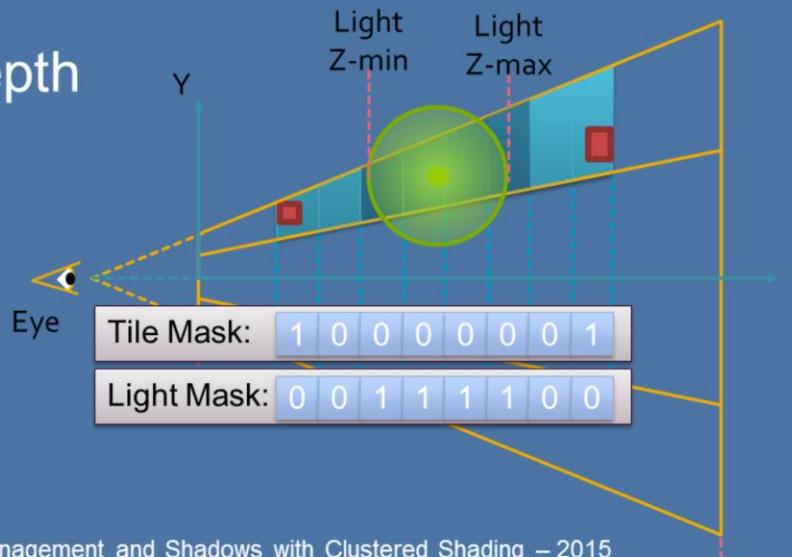


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Some techniques that attempts to address this exists
- One is 2.5D culling, which

# 2.5D culling [Harada12]

- Bit mask in depth direction
- And for light

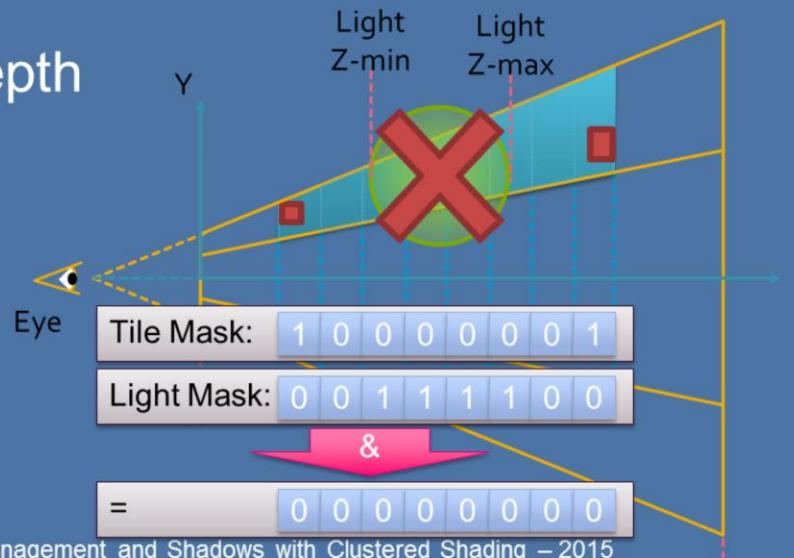


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Some techniques that attempts to address this exists
- One is 2.5D culling, which

# 2.5D culling [Harada12]

- Bit mask in depth direction
- And for light
- Binary AND
  - Discard if no overlap

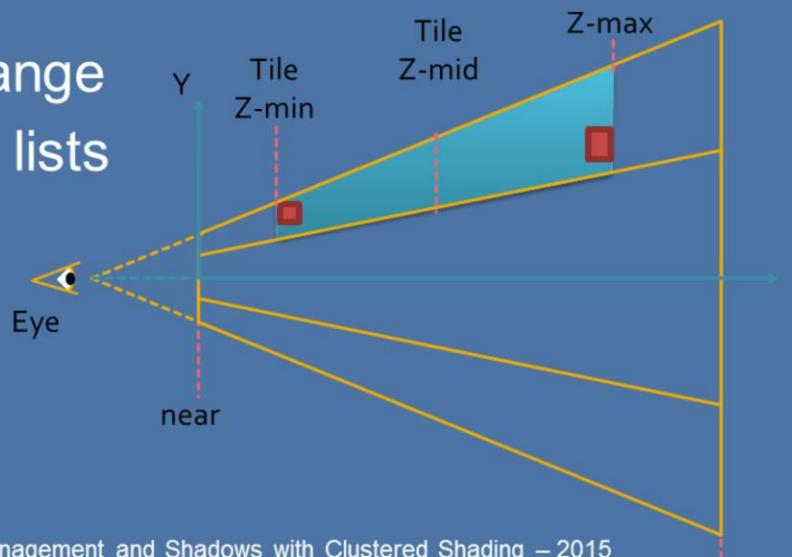


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- This takes care of some lights overlapping empty space.

## 'bimodal clusters' / 'HalfZ' [Lauritzen12][Thomas15]

- Split Depth Range
- Separate light lists

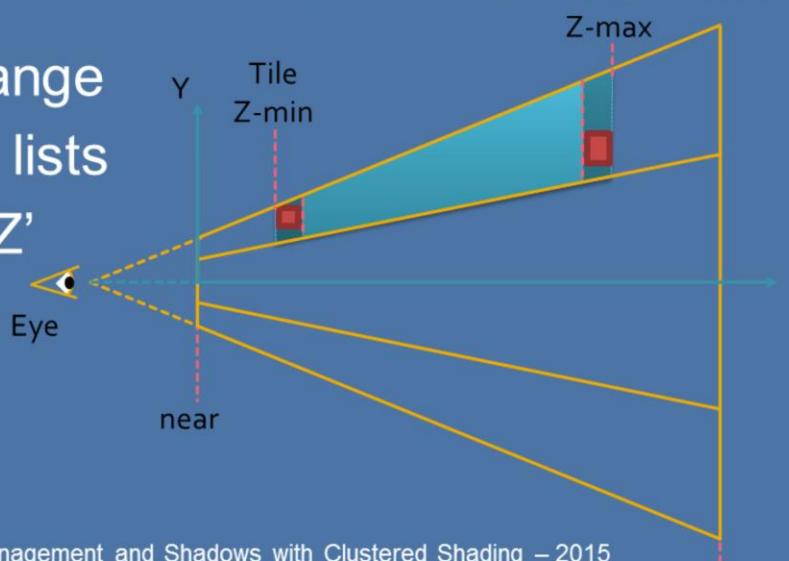


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Some techniques that attempts to address this exists
- One is 2.5D culling, which

## 'bimodal clusters' / 'HalfZ' [Lauritzen12][Thomas15]

- Split Depth Range
- Separate light lists
- 'Modified HalfZ'
- Shrink to fit



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Some techniques that attempts to address this exists
- One is 2.5D culling, which

# Tiled extensions

- Attempts to address this problem
- 2.5D culling [Harada12]
- 'bimodal clusters' / 'HalfZ' [Lauritzen12][Thomas15]
- Problems
  - Lack of generality: Slopes / multiple layers
  - No solution for transparency

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- In summary, these techniques lack generality, and don't really handle slopes.
- Also they both require a pre-existing depth buffer, and thus do not support transparency.

Clustered Deferred, Clustered Forward

# CLUSTERED SHADING

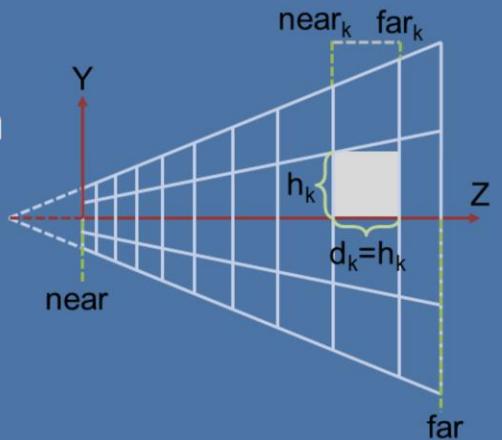
Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- This brings us to the second of the modern techniques, clustered shading.
- It is squarely aimed at addressing this problem, and making sure it stays solved.

~25 - 28

# Clustered Shading – Key Idea

- Add the 3rd dimension
  - Tile also in depth direction
- Cluster
  - Groups *related* samples

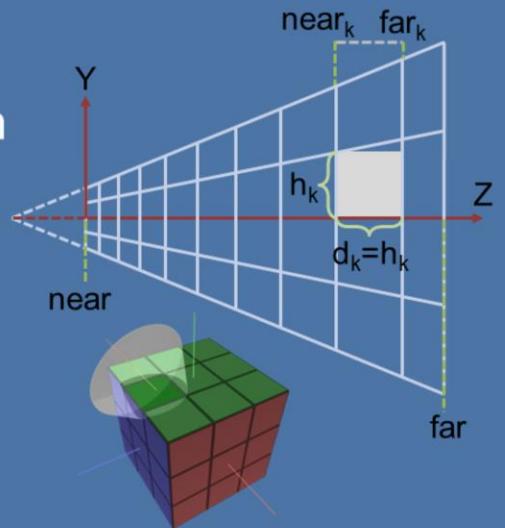


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Clustered shading can be seen as a generalization of tiled shading, and can be said to cover pretty much any way of
- grouping of related sample points.
- For what we're interested in, this relation pretty much boils down to 3D Euclidean distance.
- which means that they are nearby in space.<click>

# Clustered Shading – Key Idea

- Add the 3rd dimension
  - Tile also in depth direction
- Cluster
  - Groups *related* samples
  - Also > 3 dimensions (e.g. normals) [Olsson12a]
  - Or other attributes...

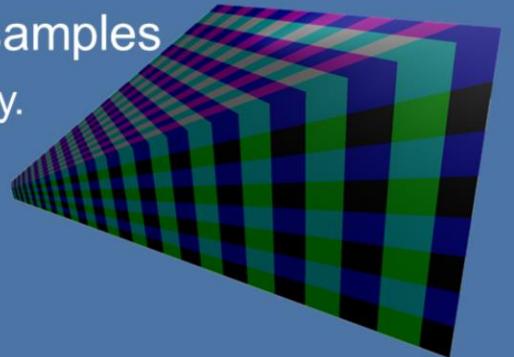


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- But we could cluster in yet higher dimensions, for example based on normal direction...<click>
- ...or on completely different attributes, like reflectance.
- But here we'll focus on 3D position, and we don't use any fancy k-means clustering algorithm (though that would be interesting)
- But simply bin the view frustum into grid cells – you may have heard the term froxels, this is pretty much the same.

# Clustered Shading Properties

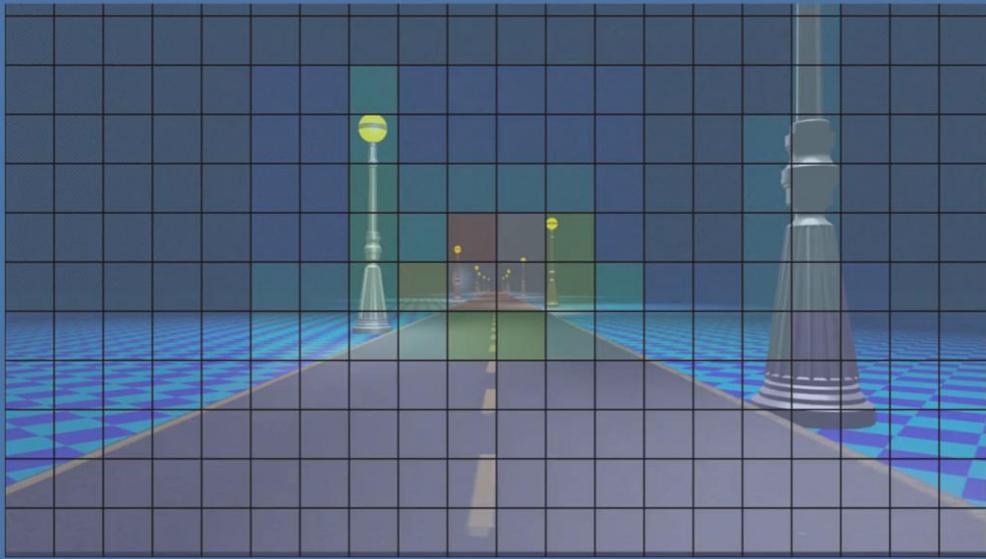
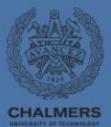
- 3D subdivisions
  - Bounded volume around samples
    - Shading cost  $\sim$  Light density.



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- So by doing 3D subdivisions we get a much more localized and consistent volume around each group of samples.
- Which enables a correlation between light density in the scene and the run time shading cost.
- This is important for control and predictability of performance.
- Lets take a look at how it works out using the road example from before.

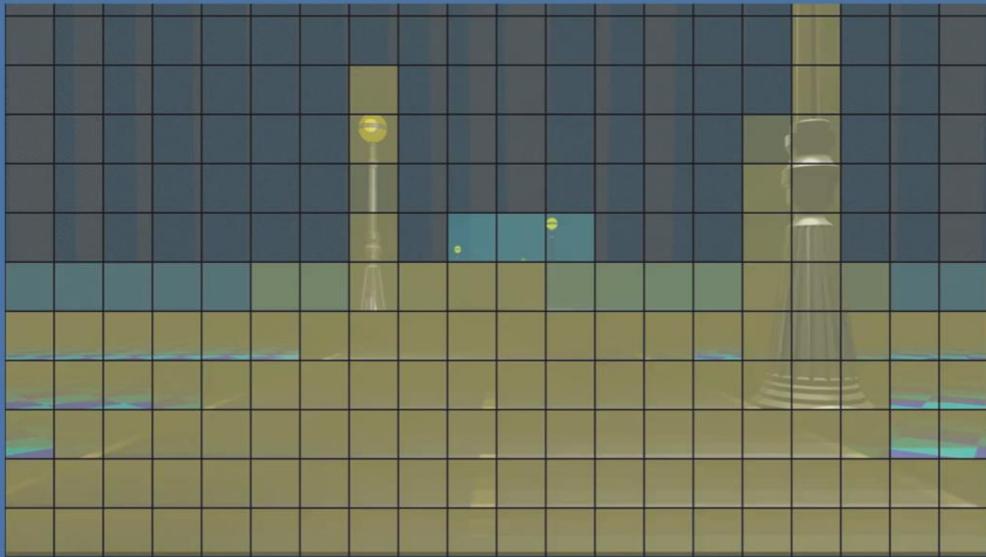
# Troublesome Tiles



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Recall the tiles.

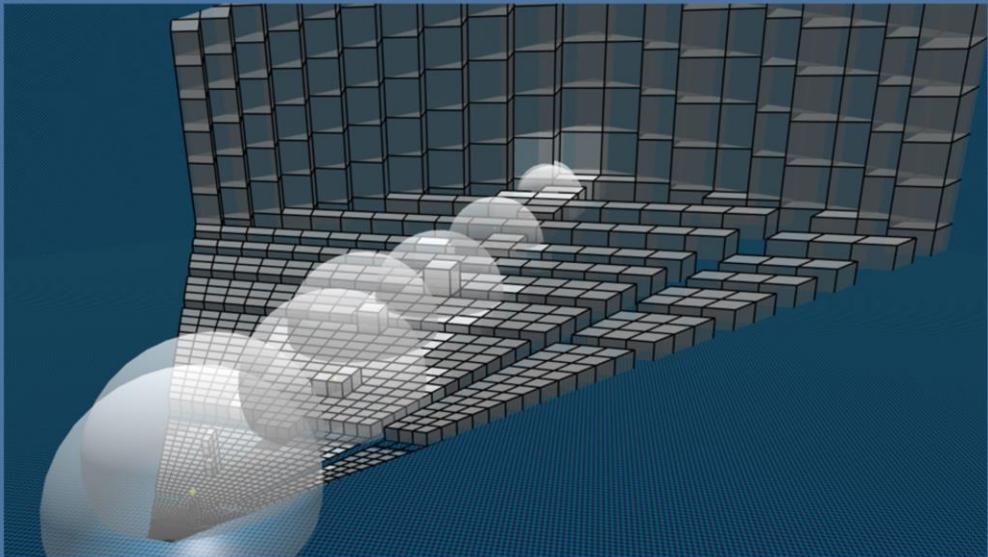
# Nice Clusters



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- And here the same view, but with clusters.
- As we can see, the clusters will overlap only a few lights each
- And follow the scene geometry well.

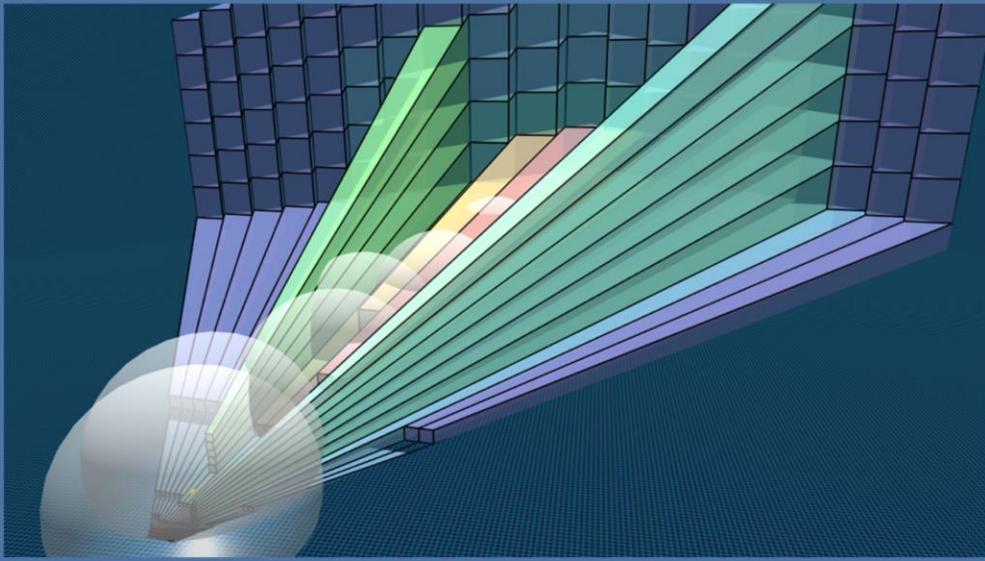
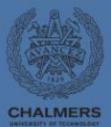
# Nice Clusters



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- And here the same view, but with clusters.
- As we can see, the clusters will overlap only a few lights each
- And follow the scene geometry well.

# Troublesome Tiles

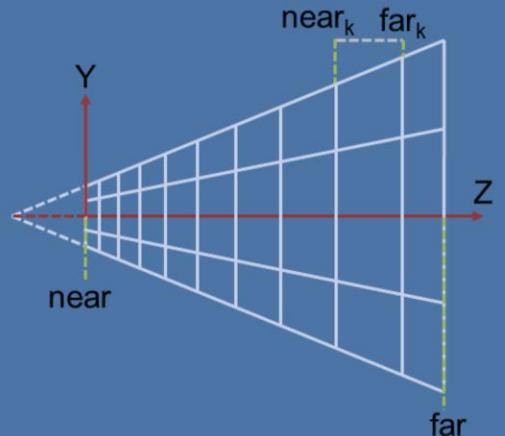


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Going back to the tiles, it looks like this...

# Cluster key / index

- Cluster key:
  - $ck = (i, j, k)$  integer tuple
- $i, j = 2D$  tile id
  - `gl_FragCoord.xy`
- $k = \approx \log(\text{view space } z)$

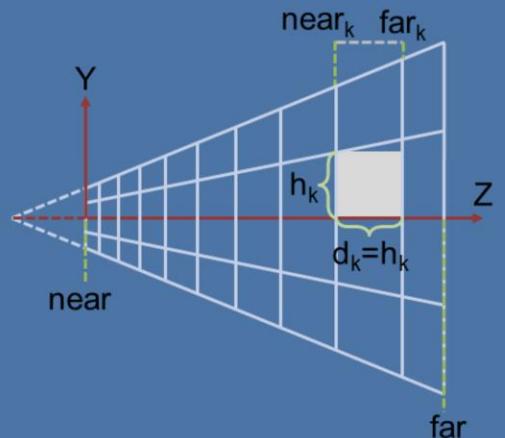


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- So, a bit more detail on how we actually build our clusters?
- The original scheme from the paper, though you will see a couple of modifications later, is to use regular subdivision starting from 2D tiles,
- which provides the  $i$  and  $j$  coordinates.
- While  $k$  is a logarithmic function of the view space  $Z$  of the sample, not simply the logarithm, for the exact equation see the paper.
- This defines a 3D grid within the view volume. <click>

# Cluster key / index

- Cluster key:
  - $ck = (i, j, k)$  integer tuple
- $i, j = 2D$  tile id
  - `gl_FragCoord.xy`
- $k = \approx \log(\text{view space } z)$

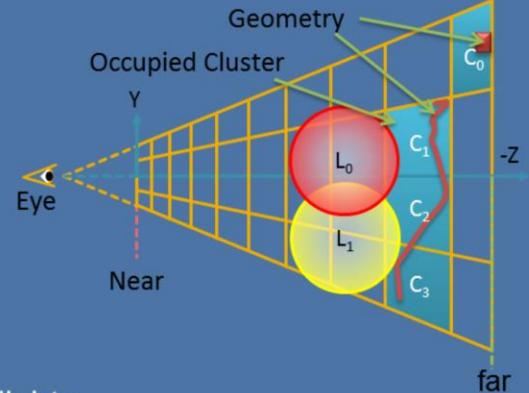


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- The logarithmic subdivision creates self similar clusters that are as cube like as possible.
- This means that as clusters become larger further away, we get a kind of LOD behavior and do not end up with insane numbers of clusters far away, for a wide range of view parameters.

# Light Assignment

- Clusters
  - Proxy for visible geometry
    - Bounds view samples
      - 3D boxes
    - Orders of magnitudes fewer
- Lights
  - Bounding spheres/cones
  - Cluster/Light Pairs
    - Links overlapping clusters and lights



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Light assignment, just like for tiled shading, is finding the overlap between lights and clusters...
- and recording this in a list per cluster.
- We will cover some concrete approaches to this later in the course.

# Variations on a Theme

- Sparse vs. Dense Cluster Grid
- Explicit vs. Implicit Clusters
- Deferred vs. Forward
- Hierarchical Clusters or not
- Supporting transparency

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- So that was the super high level overview of clustered shading.
- Now lets try to look at some of the important design decisions you can pick when implementing it.

# Sparse vs. Dense Cluster Grid

- Sparse Cluster Grid
  - Only store cells that contain samples
  - Requires Pre(-Z) pass / Deferred
  - No redundant light assignment
  - Geometry info useful for other things!
- Dense Cluster Grid
  - Must assign lights to all clusters
  - Can be done on CPU / Asynchronously
  - Can access any point in view volume
  - SAME shading cost as for sparse!

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

# Explicit vs. Implicit Clusters

- Explicit cluster bounds
  - Actual bounding box of samples
  - Some storage
  - Some cost to build
  - Tight bounds
  - Extra geometry pass for forward shading
- Implicit cluster bounds
  - Implied by grid coordinate
  - No storage
  - Can have large empty space.

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

# Deferred vs. Forward Shading

- No fundamental decision
  - Easy to change
- Deferred
  - Easier to support sparse grid
- Forward
  - MSAA, Transparency
  - Vary shading model
- Both
  - Deferred for opaque
  - Forward for transparent
  - Same or separate grids

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

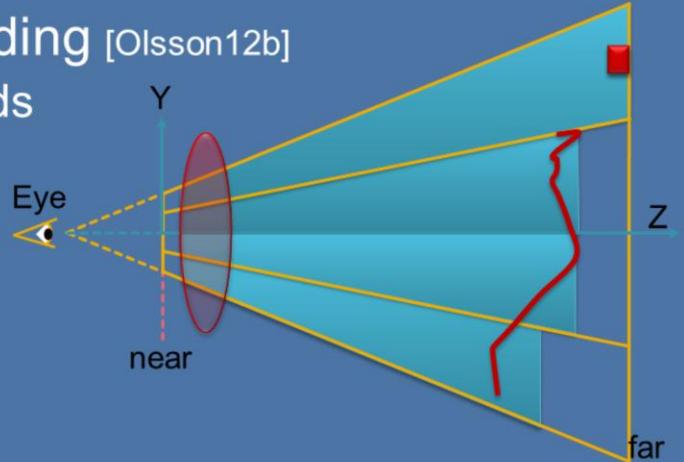
# Cluster Hierarchy

- Cheap to build [Sintorn14]
- Hierarchical representation of visible samples.
- Enables novel techniques
  - Efficient Shadow Volumes
  - Shadow maps (Talk #3)
  - Splatting in general.

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

# Transparent Geometry

- Tiled Forward Shading [Olsson12b]
  - Extend tile Z bounds
    - To Min Z
    - Or Near plane



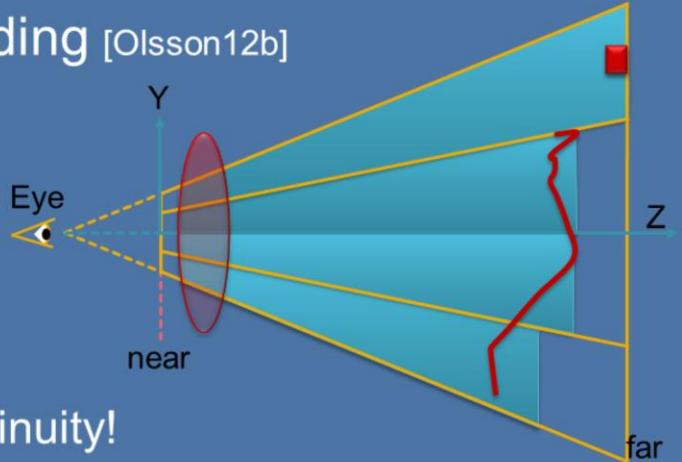
Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- It's simple to support transparent geometry with tiled forward shading.
  - For example by finding the min and max Z values, or just extending tiles to the near plane.
  - However...

# Transparent Geometry

- Tiled Forward Shading [Olsson12b]

- Problems
  - View dependent
  - Degenerates to 2D
  - Full screen discontinuity!

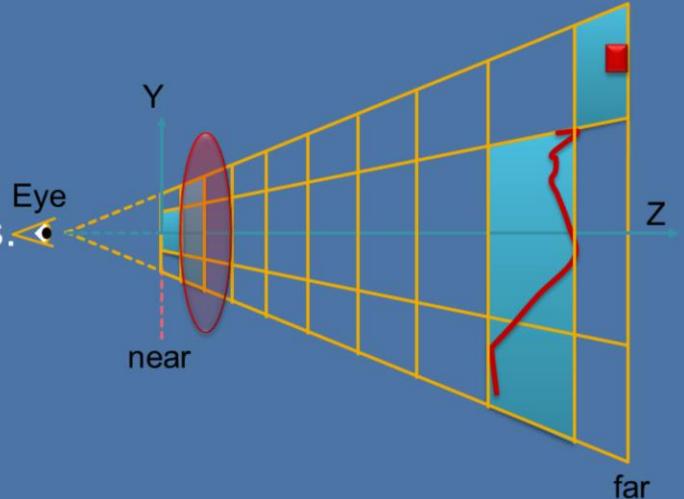


Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- The result of transparent geometry covering the view is effectively the loss of the depth range optimization
  - which can be a very bad thing™.
- Again, the biggest issue here is that it is view dependent, and so will only be possible to determine at run time
  - and, as we all know, this kind of problem usually starts showing up five minutes after we shipped a build to the publisher.

# Clustered Forward Shading

- Dense cluster grid
  - Just works!
- Sparse cluster grid
  - Pre-geometry pass.
  - Flag used clusters.
    - Side effect in fragment shader.
    - Set cell flag to one.



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

- Clustered Shading does not suffer from this problem, as each cluster represents a fixed section of 3D space.
- The only problem we face is trying to determine what clusters are need such that we can assign lights to them.
- This can be done by rendering a pre-pass with the transparent geometry.
- This can be performed after a regular G-buffer or pre-z pass, with color and depth writes turned off, and sets the used clusters to one as a side effect in the fragment shader.
- The algorithm is otherwise as before.
- Note that clustered shading provides an efficient way to solve the light assignment problem for any transparency technique, not just good old order dependent transparency.

# End!



Real-Time Many-Light Management and Shadows with Clustered Shading – 2015

# Technique Comparison

	Trad. Deferred	Tiled Deferred	Tiled Forward	Clustered Deferred	Clustered Forward
Bandwidth Use	High	Low	Low	Low	Low
Transparency	No	With Fwd →	Yes	With Fwd →	Yes
MSAA	Maybe	Maybe	Yes	Maybe	Yes
Shadow Map reuse	Yes	No	No	No	No
Geometry Passes	1	1	1–2*	1	1–2 <sup>+</sup>
Register Pressure	Low	High	High	High	High
Innermost loop	Pixels	Lights	Lights	Lights	Lights
FB Precision	High	Low	Low	Low	Low
View dependence	Low <sup>#</sup>	High	High	Low	Low
Vary Shading Models	Hard	Hard	Simple	Hard	Simple

Footnotes:

- \*: Depends on whether pre-z pass is used, which is pretty much required if the scene has any depth complexity
- +: 1 Pass is generally fine if a dense grid is used. Then a pre-z pass is only used to prime depth buffer for overdraw handling.
- #: If the stencil buffer optimization is NOT used (and perhaps it is not practical today) then view dependence might be quite high because of light overdraw!

# References

- [Persson13] Practical Clustered Deferred and Forward Shading, Persson & Olsson 2013, <http://advances.realtimerendering.com/s2013>
- [Olsson12a] Clustered Deferred and Forward Shading, Olsson et.al. 2012
- [Olsson12b] Tiled and Clustered Forward Shading, Olsson et.al. 2012
- [Valient14] Reflections and volumetrics of Killzone Shadow Fall , SIGGRAPH'14, <http://advances.realtimerendering.com/s2014>
- [Andersson14] Rendering Battlefield 4 with Mantle, GDC 2014
- [Shulz14] Moving to the Next Generation - The Rendering Technology of Ryse, GDC 2014.
- [Tebbs92] Brice Tebbs, Ulrich Neumann, John Eyles, Greg Turk, and David Ellsworth. Parallel architectures and algorithms for real-time synthesis of high quality images using deferred shading. 1992.
- [Saito90] Takafumi Saito and Tokiichiro Takahashi. Comprehensible rendering of 3-d shapes. SIGGRAPH Comput. Graph., 24(4):197–206, 1990.
- [Hargreaves04] Shawn Hargreaves and Mark Harris. Deferred shading. NVIDIA Developer Conference: 6800 Leagues Under the Sea, 2004.
- [Shishkovtsov05] Oles Shishkovtsov. Deferred shading in S.T.A.L.K.E.R. In GPU Gems 2. Addison-Wesley, 2005.
- [Arvo03] Jukka Arvo and Timo Aila. Optimized shadow mapping using the stencil buffer. JGT, 8(3):23–32, 2003.
- [Harada12] A 2.5D culling for forward+. In SIGGRAPH Asia 2012 Technical Briefs, ACM, 18:1–18.4.
- [Lauritzen12] Andrew Lauritzen, Intersecting Lights with Pixels, In SIGGRAPH 2012 Courses, Beyond Programmable Shading,
- [Ferrier11] Alex Ferrier and Christina Coffin, Deferred shading techniques using frostbite in "Battlefield 3" and "Need for Speed the Run", SIGGRAPH 2011 Talks
- [Sintorn14] Erik Sintorn , Viktor Kämpe , Ola Olsson and Ulf Assarsson, Per-Triangle Shadow Volumes Using a View-Sample Cluster Hierarchy, I3D 2014
- [Trebilco07] Damian Trebilco, 2007 (Shader X7 2009), GitHub: <https://github.com/dtrebilco/lightindexed-deferredrender>

Papers / Slides / Demo implementation with source: <http://www.cse.chalmers.se/~olaolss>

Real-Time Many-Light Management and Shadows with Clustered Shading – 2015